



THE UNIVERSITY
of EDINBURGH

Text Technologies for Data Science

INFR11145

Text Classification

Instructor:
Björn Ross

12-Nov-2025

Pre-Lecture

- Today
 - Lecture 1 (Text classification): Theory
 - Coursework 2
- Next week
 - Lecture cancelled due to strikes

Lecture Objectives

- Learn about text basics of text classification
 - Definition
 - Types
 - Methods and models

Björn Ross, TTDS 2025/2026



3

Definition

Björn Ross, TTDS 2025/2026



4

Text Classification

- **Text classification** is the process of classifying documents into predefined categories based on their content.
 - Input: Text (document, article, sentence)
 - Task: Classify into predefined one/multiple categories
 - Categories:
 - Binary: relevant/irrelevant, spam .. etc.
 - Few: sports/politics/comedy/technology
 - Hierarchical: patents

Classification is and is not

- **Classification** (a.k.a. “**categorization**”): a common technology in data science; studied within pattern recognition, statistics, and machine learning.
- Definition:
the activity of **predicting** to which among a **predefined finite set** of groups (“classes”, or “categories”) a data item belongs to
- Formulated as the task of generating a hypothesis (or “**classifier**”, or “**model**”)

$$h : D \rightarrow C$$

where $D = \{x_1, x_2, \dots\}$ is a domain of data items and $C = \{c_1, \dots, c_n\}$ is a finite set of classes (the **classification scheme**)

Classification is and is not

- Different from clustering, where the groups (“clusters”) and their number are not known in advance
- Unsuitable when class membership can be determined with certainty (relatively easily)
 - e.g., predicting whether a natural number belongs to *Prime* or *Non-Prime* is not classification
- In text classification, data items are
 - **Textual:** e.g., news articles, emails, sentences, queries, etc.
 - **Partly textual:** e.g., Web pages

Types of classification

Types of Classification

- **Binary:**

item to be classified into one of two classes

$$h : D \rightarrow C, C = \{c_1, c_2\}$$

- e.g., Spam/not spam, offensive/not offensive, rel/irrel

- **Single-Label Multi-Class (SLMC)**

item to be classified into only one of n possible classes.

$$h : D \rightarrow C, C = \{c_1, \dots, c_n\}, \text{ where } n > 2$$

- e.g., Sports/politics/entertainment, positive/negative/neutral

- **Multi-Label Multi-Class (MLMC)**

item to be classified into none, one, two, or more classes

$$h : D \rightarrow 2^C, C = \{c_1, \dots, c_n\}, \text{ where } n > 1$$

- e.g., Assigning CS articles to classes in the ACM Classification System
- Usually solved as n independent binary classification problems

Dimension of Classification

- Text classification may be performed according to several dimensions (“axes”) orthogonal to each other
- by **topic**; by far the most frequent case, its applications are global
- by **sentiment**; useful in market research, online reputation management, social science and political science
- by **language** (a.k.a. “language identification”); useful, e.g., in query processing within search engines
- by **genre**; e.g., AutomotiveNews vs. AutomotiveBlogs, useful in website classification and others;
- by **author** (a.k.a. “authorship attribution”), by native language (“native language identification”), or by gender; useful in forensics and cybersecurity
- by **usefulness**; e.g., product reviews
-

Methods and models

Björn Ross, TTDS 2025/2026



11

Methods and models

- Rule-based classification
- Supervised-learning classification
 - Traditional features
 - Word embeddings
- Pre-trained language models
 - Supervised fine-tuning for classification
 - Zero-shot classification

Björn Ross, TTDS 2025/2026



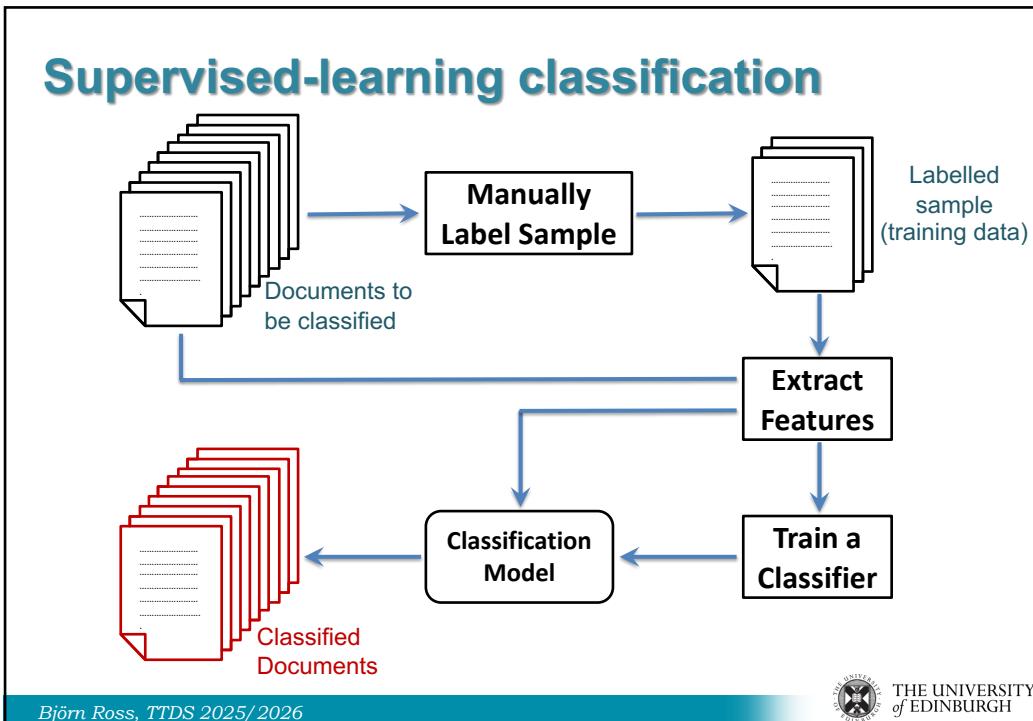
12

Rule-based classification

- An old-fashioned way to build text classifiers was via knowledge engineering, i.e., manually building classification rules
 - E.g., (Viagra or Sildenafil or Cialis) → Spam
 - E.g. (#MAGA or America great again) → support Trump
- Common type: dictionary-based classification
- Disadvantages:
 - Expensive to setup and to maintain
 - Depends on few keywords → bad coverage (recall)

Supervised-learning classification

- A generic (task-independent) learning algorithm is used to train a classifier from a set of manually classified examples
- The classifier learns, from these training examples, the characteristics a new text should have in order to be assigned to class c
- Advantages:
 - Generating training examples cheaper than writing classification rules
 - Easy update to changing conditions (e.g., addition of new classes, deletion of existing classes, shifted meaning of existing classes, etc.)



15

Extract Features

- In order to be input to a learning algorithm (or a classifier), all training (or unlabeled) documents are converted into **vectors** in a common **vector space**
- The dimensions of the vector space are called **features**
- In order to generate a vector-based representation for a set of documents D , the following steps need to be taken
 1. Feature Extraction
 2. Feature Selection or Feature Synthesis (optional)
 3. Feature Weighting

Björn Ross, TTDS 2025/2026

 THE UNIVERSITY
of EDINBURGH

16

Step 1: Feature Extraction

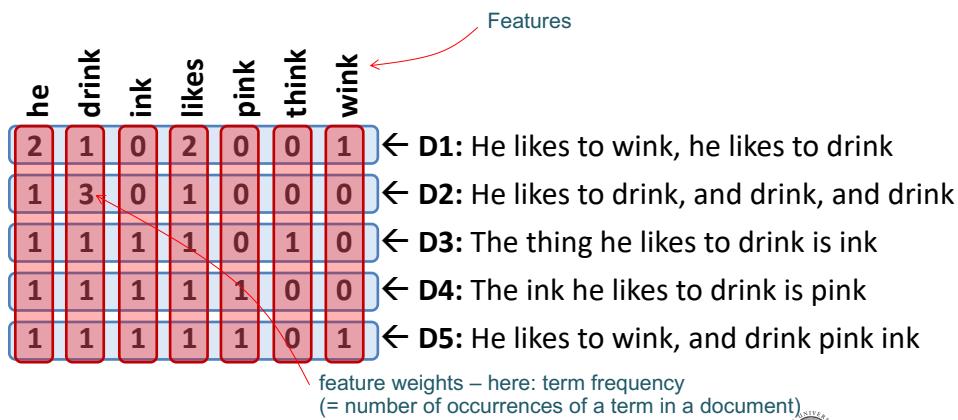
- What are the features that should be different from one class to another?
- Simplest form: Bag-of-words (BOW)
 - Each term in a document is a feature
 - Feature space size = vocabulary in all docs
 - Standard IR preprocessing steps are usually applied
 - Tokenisation, stopping, stemming

Step 1: Feature Extraction

- Bag-of-words (BOW)
 - Recall from Indexing lecture how we represented **documents** and words as vectors

Step 1: Feature Extraction

- Bag-of-words (BOW)
 - Recall from Indexing lecture how we represented documents and **words** as vectors



Björn Ross, TTDS 2025/2026

THE UNIVERSITY
of EDINBURGH

19

Step 1: Feature Extraction

- Other simple features forms:
 - Word n-grams (bigrams, trigrams,)
 - Much larger + more sparse
 - Sometimes char n-grams are used
 - Especially for degraded text (OCR or ASR outputs)

Björn Ross, TTDS 2025/2026

THE UNIVERSITY
of EDINBURGH

20

Step 1: Feature Extraction

- What other text features could be used?
- Sentence structure:
 - POS (part-of-speech tags)
 - Syntactic tree structure
- Topic-based features:
 - LDA topics
 - NEs (named entities) in text
 - Links / Linked terms
- Non-textual features:
 - Average doc\sentence\word length
 - % of words start with upper-case letter
 - % of links/hashtags/emojis in text

Step 1: Feature Extraction

- What preprocessing to apply?
 - Case-folding? `really` vs `Really` vs `REALLY`
 - Punctuation? “?”, “!”, “@”, “#”
 - Stopping? “he”, “she”, “what”, “but”
 - Stemming? “`replaced`” vs “`replacement`”
- Other Features:
 - Starts with capital letter, all caps
 - Repeated characters “`congraaaaats`” “`help!!!!!!`”
 - Scores from dictionaries and lexicons (e.g. LIWC)
- Which to choose?
 - Classification task/application

Step 2: Feature Selection

- Number of distinctive features = length of feature vector
- Vector can be of length in the order of 10^6 , and might be sparse
 - High computational cost
 - Overfitting
- What are the most important features among those?
 - e.g. Reduce from 10^6 to 10^4
- For each class, find the top representative k features for it → get the Union over all classes → reduced feature space

Step 2: Feature Selection Functions

- Document frequency
 - % of docs in class c_i that contain the term t_k
 - Very basic measure. Will select stop words as features

$$\#(t_k, c_i) = P(t_k | c_i)$$
- Mutual Information
 - How much we learn from the presence or absence of term t_k about whether or not a document is in class c_i
 - Often used in feature selection in text classification

$$MI(t_k, c_i) = \sum_{c \in \{c_i, c_i\}} \sum_{t \in \{t_k, t_k\}} P(t, c) \cdot \log_2 \frac{P(t, c)}{P(t) \cdot P(c)}$$
- Pearson's Chi-squared (χ^2)
 - used more in comparisons between classes

Step 2: Feature Selection Functions

Function	Denoted by	Mathematical form
Document frequency	$\#(t_k, c_i)$	$P(t_k c_i)$
DIA association factor	$z(t_k, c_i)$	$P(c_i t_k)$
Information gain	$IG(t_k, c_i)$	$\sum_{c \in \{c_i, \bar{c}_i\}} \sum_{t \in \{t_k, \bar{t}_k\}} P(t, c) \cdot \log \frac{P(t, c)}{P(t) \cdot P(c)}$
Mutual information	$MI(t_k, c_i)$	$\log \frac{P(t_k, c_i)}{P(t_k) \cdot P(c_i)}$
Chi-square	$\chi^2(t_k, c_i)$	$\frac{ Tr \cdot [P(t_k, c_i) \cdot P(\bar{t}_k, \bar{c}_i) - P(t_k, \bar{c}_i) \cdot P(\bar{t}_k, c_i)]^2}{P(t_k) \cdot P(\bar{t}_k) \cdot P(c_i) \cdot P(\bar{c}_i)}$
NGL coefficient	$NGL(t_k, c_i)$	$\sqrt{\frac{ Tr \cdot [P(t_k, c_i) \cdot P(\bar{t}_k, \bar{c}_i) - P(t_k, \bar{c}_i) \cdot P(\bar{t}_k, c_i)]}{P(t_k) \cdot P(\bar{t}_k) \cdot P(c_i) \cdot P(\bar{c}_i)}}$
Relevancy score	$RS(t_k, c_i)$	$\log \frac{P(t_k c_i) + d}{P(\bar{t}_k \bar{c}_i) + d}$
Odds Ratio	$OR(t_k, c_i)$	$\frac{P(t_k c_i) \cdot (1 - P(t_k \bar{c}_i))}{(1 - P(t_k c_i)) \cdot P(t_k \bar{c}_i)}$
GSS coefficient	$GSS(t_k, c_i)$	$P(t_k, c_i) \cdot P(\bar{t}_k, \bar{c}_i) - P(t_k, \bar{c}_i) \cdot P(\bar{t}_k, c_i)$

Björn Ross, TTDS 2025/2026



25

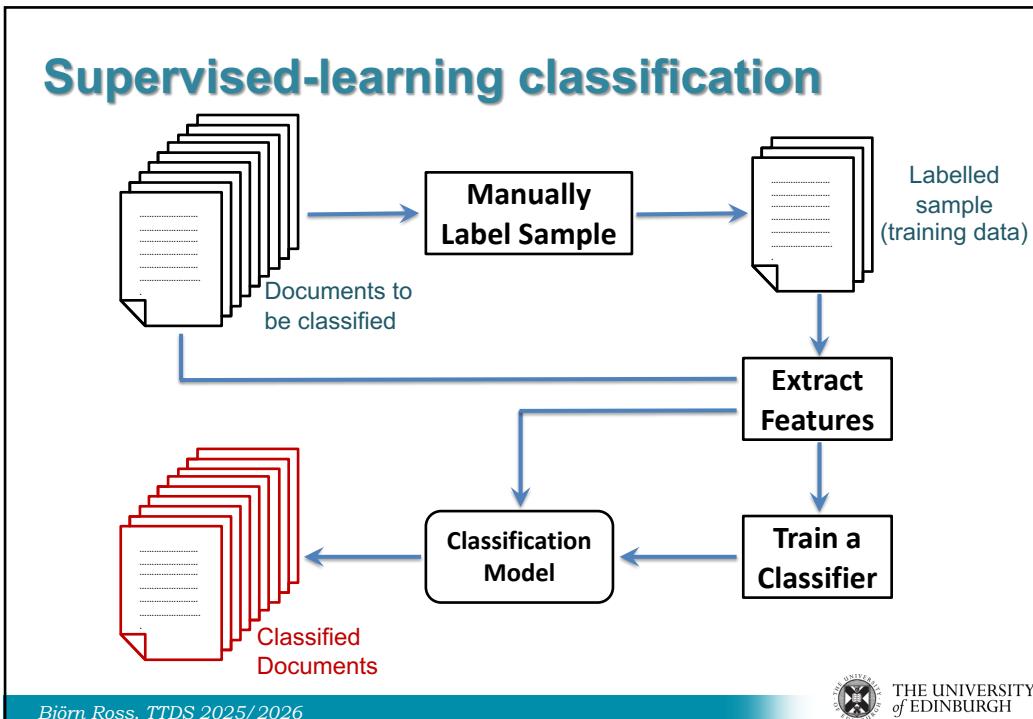
Step 3: Feature Weighting

- Attributing a value to feature t_k in document d_i
This value may be
 - binary** (representing presence/absence of t_k in d_i);
 - numeric** (representing the importance of t_k for d_i);
obtained via feature weighting functions in the following two classes:
 - unsupervised**: e.g., tfidf or BM25,
 - supervised**: e.g., $tf^* MI$, $tf^* \chi^2$
- Similarity between two vectors may be computed e.g.
via **cosine similarity**
- Scaling** can be important!

Björn Ross, TTDS 2025/2026



28



29

Training a Classifier

- For **binary** classification, essentially any supervised learning algorithm can be used for training a classifier; classical choices include
 - Support vector machines (SVMs)
 - Random forests
 - Naïve Bayesian methods
 - Lazy learning methods (e.g., k-NN)
 - Logistic Regression
 -
- The “**No-free-lunch principle**” (Wolpert, 1996) → *there is no learning algorithm that can outperform all others in all contexts*
- Implementations need to cater for
 - the very high dimensionality
 - the sparse nature of the representations involved

Björn Ross, TTDS 2025/2026

THE UNIVERSITY
of EDINBURGH

30

Training a Classifier

- For **Multiclass classification**, some learning algorithms for binary classification are “SLMC-ready”; e.g.
 - Decision trees
 - Random forests
 - Naive Bayesian methods
 - Lazy learning methods (e.g., k-NN)
 - Neural networks
- For other learners (notably: SVMs) to be used for SLMC classification, combinations / cascades of the binary versions need to be used
 - e.g. multi-class classification SVM
 - Could be directly used for MLMC as well

Word embeddings

- More complex representation of words as vectors
 Recall the **term vectors** in traditional indexing or classification, using term frequency as weights:
 - Sparse (most values are 0)
 - Capture semantics only incidentally (similar vectors are terms that appear together)

he	drink	ink	likes	pink	think	wink	
2	1	0	2	0	0	1	← D1: He likes to wink, he likes to drink
1	3	0	1	0	0	0	← D2: He likes to drink, and drink, and drink
1	1	1	1	0	1	0	← D3: The thing he likes to drink is ink
1	1	1	1	1	0	0	← D4: The ink he likes to drink is pink
1	1	1	1	1	0	1	← D5: He likes to wink, and drink pink ink

Word embeddings

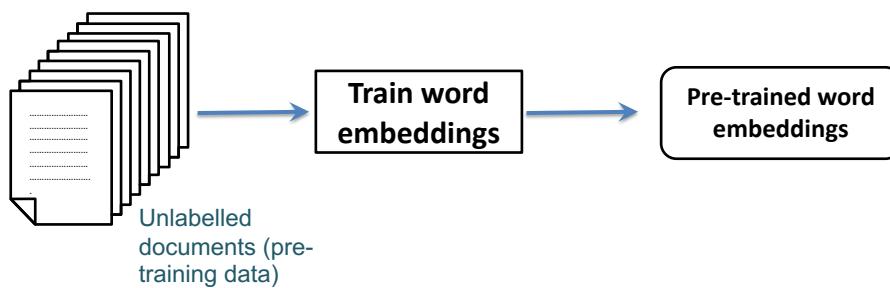
- More complex representation of words as vectors
 - Dense (all entries are non-zero)
 - Capture semantics (similar words have similar vectors)

he	drink	ink
0.123	0.521	0.313
0.451	0.987	0.812
0.938	0.141	0.411
...

(many dimensions e.g. 300)

Word embeddings

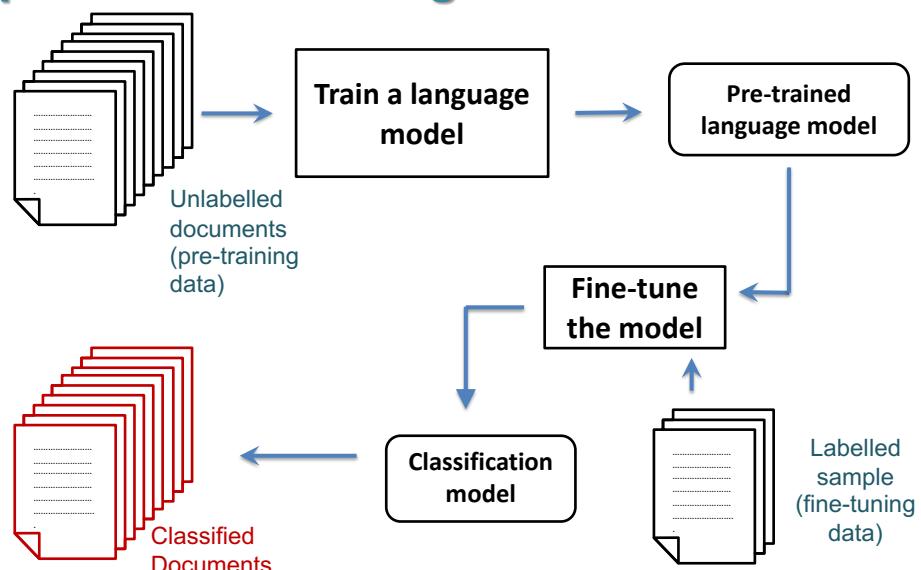
- Obtained through self-supervised learning to learn relationships between words
 - Predict a centre word given surrounding context words (CBOW)
 - Predict context words given target word (skip-gram)
- Can be done on a large unlabelled pre-training corpus
- Helps with out-of-vocabulary problems



Pre-trained language models

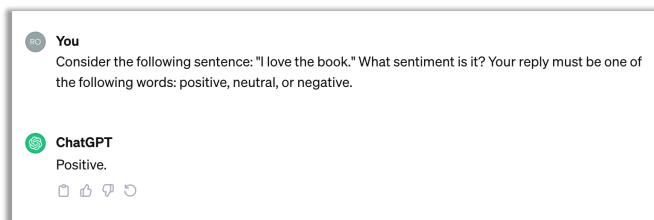
- Obtained through self-supervised learning
 - Predict the next word: The Queen of [...] (*next token prediction*)
 - Predict a masked word: The [...] of England. (*masked language modelling*)
 - ...
- Also done on large unlabelled pre-training corpora
- Penultimate layer of network can be used to generate **contextualised word embeddings** for other language-based tasks
- Basis for many state-of-the-art text classifiers
 - e.g. BERT (DistilBERT, RoBERTa..), XLNet, etc.

Supervised fine-tuning



Zero-shot classification

- Using language models for classification directly without giving training examples (= skipping the fine-tuning step)
 - A modern approach that can work well and requires little human effort
 - Depends highly on black-box models, limited opportunities for customisation and error analysis (but this is an active research area)
 - Often used as a baseline for performance comparisons

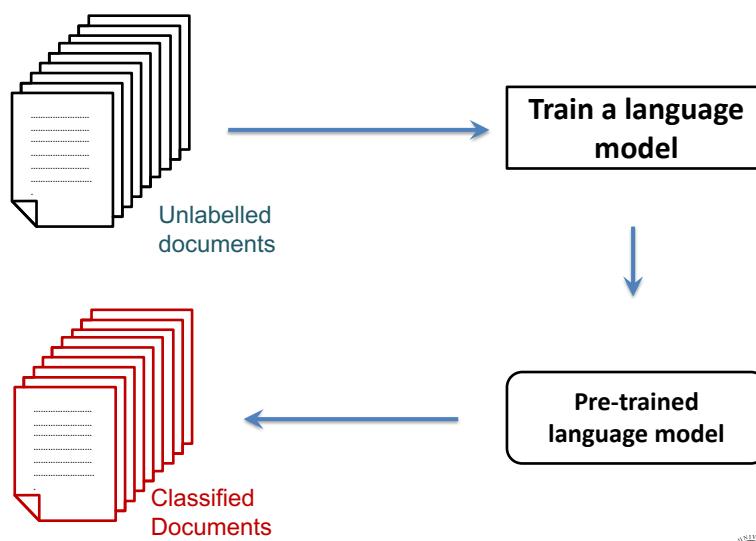


Björn Ross, TTDS 2025/2026

 THE UNIVERSITY
of EDINBURGH

37

Zero-shot classification



Björn Ross, TTDS 2025/2026

 THE UNIVERSITY
of EDINBURGH

38

Few-shot classification

- Variation of zero-shot classification: A few examples are given in the prompt
- Also skips task-specific fine-tuning
- Leverages models' capability at **in-context learning** (adapt to a task from examples in the input prompt)

Task: Identify the stance toward carbon tax in the given sentence.
Your reply must be one of: support, oppose, neutral.

Examples:

Sentence: "I think a carbon tax is necessary to fight climate change."
Target: carbon tax → support

Sentence: "Renewable energy is amazing, though I don't have strong feelings about a carbon tax."
Target: carbon tax → neutral

Now classify this sentence:
Sentence: "I support policies to reduce emissions, but I'm worried about how a carbon tax will affect me personally."
Target: carbon tax →

Target: carbon tax → oppose

Björn Ross, TTDS 2025/2026



39

Evaluation

Björn Ross, TTDS 2025/2026



40

Evaluation

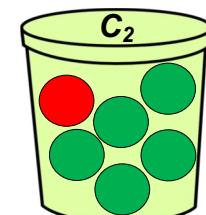
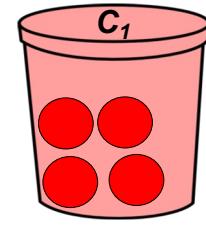
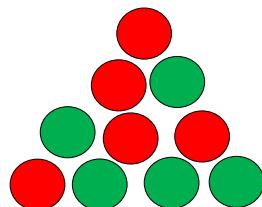
- Effectiveness (e.g. accuracy, precision, recall, F1):
 - Global effectiveness measures
 - Per class effectiveness measures
- Efficiency:
 - Speed in learning
 - SVM with linear kernel is known to be fast
 - DNNs are known to be much slower (specially with large # layers)
 - Speed in classification
 - K-NNs are known to be one of the slowest
 - Speed in feature extraction
 - BOW vs POS vs Link analysis features
- Importance of baselines

Evaluation: Baselines

- There are standard methods for creating baselines in text classification to compare your classifier with
- Most popular/simplest baselines
 - Random classification
 - Classes are assigned randomly
 - How much better is the classifier doing than random?
 - Majority class baseline
 - Assign all elements to the class that appears the most
 - How much better you are doing than if you always picked the same thing output regardless of input?
 - Simple algorithm, e.g. BOW
 - Usually used when you introduce new interesting features
 - LLMs zero-shot, e.g. GPT-4o
 - Can sometimes outperform fine-tuned models

Evaluation: Binary Classification

- Accuracy:
 - How many of the samples are classified correctly?
 - $A = (4+5)/10 = 0.9$



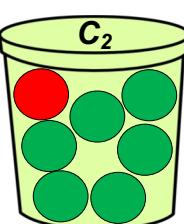
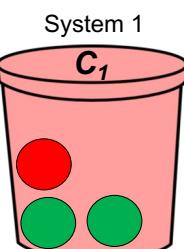
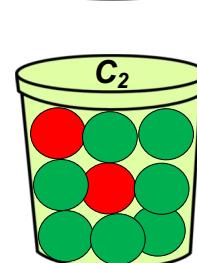
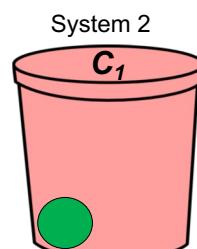
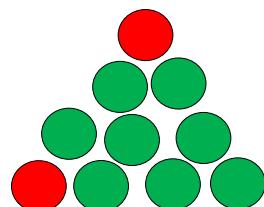
Björn Ross, TTDS 2025/2026

 THE UNIVERSITY
of EDINBURGH

43

Evaluation: Binary Classification

- $A = (1+6)/10 = 0.7$ System 1
- $A = (0+7)/10 = 0.7$ System 2
- When classes are highly unbalanced
 - Precision/recall/F1 for the rare class
 - e.g. Spam classification (detection)



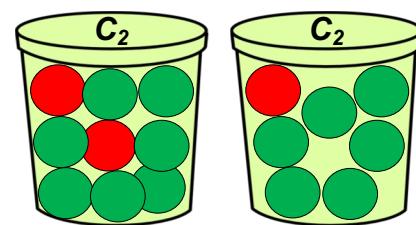
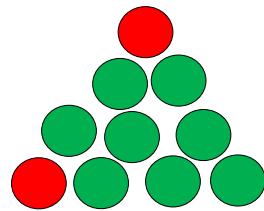
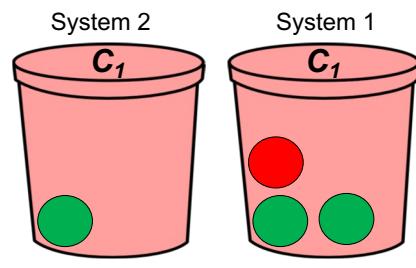
Björn Ross, TTDS 2025/2026

 THE UNIVERSITY
of EDINBURGH

44

Evaluation: Binary Classification

	System 1	System 2
Precision	$1/3 = 0.33$	$0/1 = 0$
Recall	$1/2 = 0.5$	$0/2 = 0$
F1	0.4	0



Björn Ross, TTDS 2025/2026

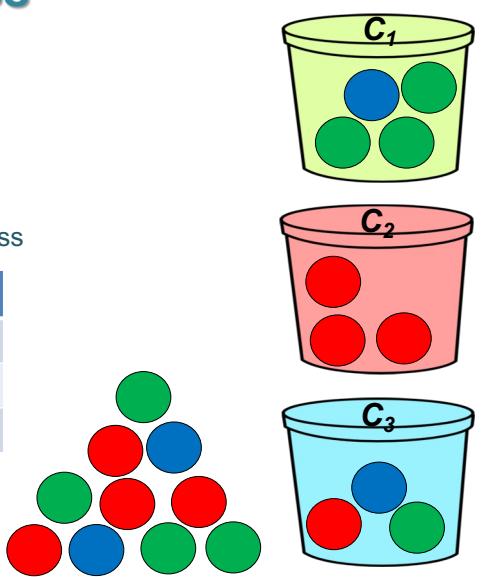
 THE UNIVERSITY
of EDINBURGH

45

Evaluation: Multi-class

- Accuracy = $(3+3+1)/10 = 0.7$
- Good measure when
 - Classes are nearly balanced
- Preferred:
 - Precision/recall/F1 for each class

	Green	Red	Blue
P	0.75	1	0.333
R	0.75	0.75	0.5
F1	0.75	0.86	0.4



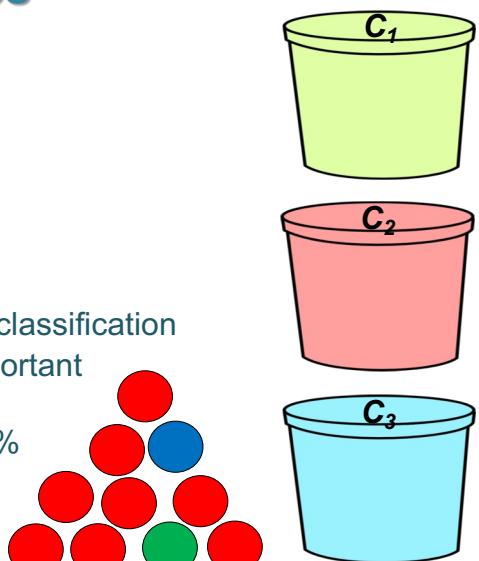
Björn Ross, TTDS 2025/2026

 THE UNIVERSITY
of EDINBURGH

46

Evaluation: Multi-class

- Majority class baseline
 - Accuracy = 0.8
 - Macro-F1 = 0.296
- Macro-F1:
 - Should be used in binary classification when two classes are important
 - e.g.: males/females while distribution is 80/20%



Björn Ross, TTDS 2025/2026

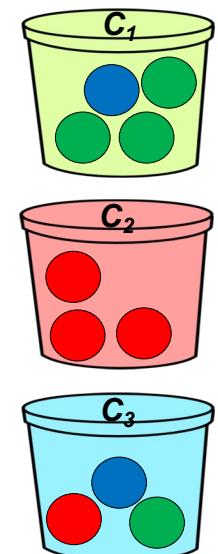
THE UNIVERSITY
of EDINBURGH

47

Error Analysis

- Confusion Matrix

		Predicted class		
		Green	Red	Blue
Actual class	Green	3	0	1
	Red	0	3	1
	Blue	1	0	1



Björn Ross, TTDS 2025/2026

THE UNIVERSITY
of EDINBURGH

48

Data splitting

- It's important to avoid overfitting
- Labelled data could be split into **two parts**
 - **Training**: used to train the classifier (e.g. **80%** of the data)
 - **Test**: used to test the performance of the trained classifier on unseen data (e.g. **20%** of the data)

Hyperparameter optimisation

- Most classifiers have some hyperparameters to be optimized
 - The C parameter in soft-margin SVMs
 - The r, d parameters of non-linear kernels
 - Decision threshold for binary SVM
- We may also try different models (SVM, Fine-tuned RoBERTa, GPT-4o zero-shot..) so we could overfit to this choice
- Usually labelled data is split into **three parts**
 - **Training**: used for training / fine-tuning (typically **80%** of the data)
 - **Development**: used to optimise hyperparameters. Apply the classifier on this data with different values of the hyperparameters and report the one that achieves the highest results (usually **10%** of the data)
 - **Test**: used to test the performance of the trained classifier with the optimal hyperparameters on these unseen data (usually **10%** of the data)
- Optimising the hyperparameters on test data is cheating!

Summary

- Text Classification tasks
- Types of text classification
- Models and methods for text classification
 - Rule-based
 - Supervised learning-based
 - Pre-trained language models
- Baselines and evaluation

Resources

- *Fabrizio Sebastiani*
Machine Learning in Automated Text Categorization
ACM Computing Surveys, 2002
Link: <https://arxiv.org/pdf/cs/0110053>
- *Yoav Goldberg*
A Primer on Neural Network Models for Natural Language Processing
Link: <https://arxiv.org/abs/1510.00726>
- *Brown, T., Mann, B., Ryder, N., Subbiah, M., Kaplan, J. D., Dhariwal, P., ... & Amodei, D. (2020). Language models are few-shot learners.* Advances in Neural Information Processing Systems, 33, 1877-1901.