

Ranked Retrieval 1: 从 match/no match 到排序分数

图文讲义版：用原 PPT 作为视觉锚点，按“概念故事线 + 直觉解释 + 考试速记”整理。

Boolean retrieval limitation → top-k ranked output → Jaccard baseline → TF/DF/IDF → TF-IDF weighting → Vector Space Model → length normalization → cosine similarity → SMART notation

10/7/25

Boolean Retrieval

- Thus far, our queries have all been Boolean.
 - Documents either: “match” or “no match”.
- Good for expert users with precise understanding of their needs and the collection.
 - Patent search uses sophisticated sets of Boolean queries and check hundreds of search results
(car OR vehicle) AND (motor OR engine) AND NOT (cooler)
- Not good for the majority of users.
 - Most incapable of writing Boolean queries.
 - Most don't want to go through 1000s of results.
 - This is particularly true for web search
 - Question: What is the most unused web-search feature?

Walid Magdy, TTDS 2025/2026



3

Ranked Retrieval

- Typical queries: free text queries
- Results are “ranked” with respect to a query
- Large result sets are not an issue
 - We just show the top k (≈ 10) results
 - We don't overwhelm the user
- Criteria:
 - Top ranked documents are the most likely to satisfy user's query
 - Score is based on how well documents match a query
Score(d,q)

Walid Magdy, TTDS 2025/2026



4

2

这节课一句话

这节课把检索从 Boolean 的“进不进名单”，升级为 ranked retrieval 的“谁最应该排在前面”。

1. Boolean retrieval 的天花板

10/7/25

Boolean Retrieval

- Thus far, our queries have all been Boolean.
 - Documents either: “match” or “no match”.
- Good for expert users with precise understanding of their needs and the collection.
 - Patent search uses sophisticated sets of Boolean queries and check hundreds of search results
(car OR vehicle) AND (motor OR engine) AND NOT (cooler)
- Not good for the majority of users.
 - Most incapable of writing Boolean queries.
 - Most don't want to go through 1000s of results.
 - This is particularly true for web search
 - Question: What is the most unused web-search feature?

Walid Magdy, TTDS 2025/2026

THE UNIVERSITY of EDINBURGH

3

Ranked Retrieval

- Typical queries: free text queries
- Results are “ranked” with respect to a query
- Large result sets are not an issue
 - We just show the top k (≈ 10) results
 - We don't overwhelm the user
- Criteria:
 - Top ranked documents are the most likely to satisfy user's query
 - Score is based on how well documents match a query
Score(d,q)

Walid Magdy, TTDS 2025/2026

THE UNIVERSITY of EDINBURGH

4

2

原 PPT 第 2 页

人话直觉

Boolean 像门口保安：只说“能进/不能进”；但搜索引擎需要前台排序：谁最值得先见。

精确定义 / 机制： Boolean retrieval 给每个 document match/no match，适合专家用户和精确约束，但普通用户通常输入 free text query，并期待系统返回 top-k ranked results。ranked retrieval 给每个 document 一个 score，再按分数排序。

考试问法：若问 ranked vs Boolean，核心区别是 binary decision vs score-based ranking/top-k。

2. Ranking output: $\text{score}(d,q)$

10/7/25

Old Example

- Find documents matching query {ink wink}
 - Load inverted lists for each query word
 - Merge two postings lists $\rightarrow$ **Linear merge**
- Apply function for matches
 - Boolean: exist / not exist = 0 or 1
 - Ranked: $f(tf, df, length, \dots) = 0 \rightarrow 1$

Matches

ink	$\Rightarrow$	3:1	4:1	5:1
wink	$\Rightarrow$	1:1	5:1	

1: $f(0,1)$
3: $f(1,0)$
4: $f(1,0)$
5: $f(1,1)$

Walid Magdy, TTDS 2025/2026

Function example: Jaccard coefficient

- a commonly used measure of overlap of two sets A and B
- $jaccard(A,B) = \frac{|A \cap B|}{|A \cup B|}$
- $jaccard(A,A) = 1$
- $jaccard(A,B) = 0$, if $A \cap B = 0$
- Example:
 - $D1 \cup D2 = \{\text{he, likes, to, wink, and, drink}\}$
 - $D1 \cap D2 = \{\text{he, likes, to, drink}\}$
 - $jaccard(D1,D2) = \frac{4}{6} = 0.6667$

D1: He likes to wink, he likes to drink
D2: He likes to drink, and drink, and drink

Walid Magdy, TTDS 2025/2026

3

原 PPT 第 3 页

人话直觉

不是把所有候选文档倒在桌上，而是给每份卷子打分，从高到低发给用户。

精确定义 / 机制：ranked retrieval 通常输出 query_id、document_id 和 score。 $\text{score}(d,q)$ 表示 document d 与 query q 的匹配程度或相关性估计。系统可以只返回 top-k，因为用户一般只看前几条。

$$d_1 > d_2 \quad \text{if} \quad \text{score}(q, d_1) > \text{score}(q, d_2)$$

考试问法：答 ranked retrieval pipeline 时要出现 score、sort、top-k 三个词。

3. Jaccard: 一个太朴素的 overlap 分数

10/7/25

Jaccard coefficient: Issues


- Does not consider **term frequency** (how many times a term occurs in a document)
- It treats all terms equally!
 - How about **rare terms** in a collection? more informative than frequent terms.
 - *He likes to drink*, shall "to" == "drink"?
- Needs more sophisticated way of **length** normalization
 - $|D1| = 3, |D2| = 1000!$
 - $D1 \rightarrow Q, D2 \rightarrow D$

Walid Magdy, TTDS 2025/2026

THE UNIVERSITY of EDINBURGH

Should terms be treated the same?

- Collection of 5 documents (balls = terms)
- Query   
- Which is the least relevant document?
- Which is the most relevant document?



Walid Magdy, TTDS 2025/2026

THE UNIVERSITY of EDINBURGH

原 PPT 第 4 页

人话直觉

Jaccard 只数两袋词有多少交集，不管“the”和“serendipity”谁更珍贵，也不管某个词出现多少次。

精确定义 / 机制：Jaccard coefficient 衡量两个集合 overlap，取 intersection 除以 union。它能作为简单 similarity，但不考虑 term frequency，所有 terms 权重相同，也缺少良好的 length normalization。

$$J(A, B) = \frac{|A \cap B|}{|A \cup B|}$$

考试问法：若问 Jaccard limitations，答不考虑 tf、不区分 rare/common terms、length normalization 弱。

4. TF: 一个词在本文档里有多响

10/7/25

TFIDF

- **TFIDF**:
Term Frequency, Inverse Document Frequency
- **$tf(t,d)$** :
number of times term t appeared in document d
 - As $tf(t,d) \uparrow \rightarrow$ importance of t in $d \uparrow \uparrow$
 - Document about IR, contains "retrieval" more than others
- **$df(t)$** :
number of documents term t appeared in
 - As $df(t) \uparrow \rightarrow$ importance if t in a collection $\downarrow \downarrow$
 - "the" appears in many document $\rightarrow$ not important
 - "FT" is not important word in financial times articles

Walid Magdy, TTDS 2025/2026

THE UNIVERSITY of EDINBURGH

9

DF, CF, & IDF

- **DF $\neq$ CF** (collection frequency)
 - $cf(t)$ = total number of occurrences of term t in a collection
 - $df(t) \leq N$ (N : number of documents in a collection)
 - $cf(t)$ can be $\geq N$
- **DF** is more commonly used in IR than **CF**
 - **CF** is still used
- **$idf(t)$** : inverse of $df(t)$
 - As $idf(t) \uparrow \rightarrow$ rare term $\rightarrow$ importance $\uparrow \uparrow$
 - $idf(t) \rightarrow$ measure of the informativeness of t

Walid Magdy, TTDS 2025/2026

THE UNIVERSITY of EDINBURGH

10

5

原 PPT 第 5 页

人话直觉

如果一篇文章反复提到 indexing，它大概率真的在讲 indexing；但第 10 次出现没有第 1 次那么惊艳。

精确定义 / 机制：term frequency $tf(t,d)$ 是 term t 在 document d 中出现次数。TF-IDF 常用 log-scaled tf ，让第一次出现贡献很大，后续出现仍加分但边际收益递减。

$$tf\text{-weight}(t,d) = 1 + \log_{10} tf(t,d)$$

考试问法：若问为什么用 log TF，答 dampen repeated occurrences，避免重复词无限加分。

5. DF/IDF: 越常见越不稀罕

10/7/25

DF vs CF

	he	drink	ink	likes	pink	think	wink	
D1: He likes to wink, he likes to drink	2	1	0	2	0	0	1	← D1: He likes to wink, he likes to drink
D2: He likes to drink, and drink, and drink	1	3	0	1	0	0	0	← D2: He likes to drink, and drink, and drink
D3: The thing he likes to drink is ink	1	1	1	1	0	1	0	← D3: The thing he likes to drink is ink
D4: The ink he likes to drink is pink	1	1	1	1	1	0	0	← D4: The ink he likes to drink is pink
D5: He likes to wink, and drink pink ink	1	1	1	1	1	0	1	← D5: He likes to wink, and drink pink ink
DF	5	5	3	5	2	1	2	
CF	6	7	3	6	2	1	2	

Walid Magdy, TTDS 2025/2026

THE UNIVERSITY of EDINBURGH

11

IDF: formula

$$idf(t) = \log_{10} \left(\frac{N}{df(t)} \right)$$

- Log scale used to dampen the effect of IDF
- Suppose $N = 1$ million →

term	df(t)	idf(t)
calpurnia	1	6
animal	100	4
sky	1,000	3
fly	10,000	2
under	100,000	1
the	1,000,000	0

Walid Magdy, TTDS 2025/2026

THE UNIVERSITY of EDINBURGH

12

6

原 PPT 第 6 页

人话直觉

“the”像空气，哪里都有，不能帮你定位；稀有词像门牌号，越少见越能指路。

精确定义 / 机制：df(t) 是包含 term t 的 document 数量，cf(t) 是 term 在 collection 中出现总次数。inverse document frequency 用 $N/df(t)$ 的 log 衡量词的区分能力：出现在越多文档中的词 IDF 越低。

$$idf(t) = \log_{10} \frac{N}{df(t)}$$

考试问法：考试常让区分 df 与 cf：df 数 documents，cf 数 occurrences。

6. TF-IDF: 本文档高频 × 全集合稀有

10/7/25

TFIDF term weighting

- One the best known term weights schemes in IR
 - Increases with the number of occurrences within a document
 - Increases with the rarity of the term in the collection
- Combines TF and IDF to find the weight of terms

$$w_{t,d} = (1 + \log_{10}tf(t,d)) \times \log_{10}(\frac{N}{df(t)})$$

- For a query q and document d , retrieval score $f(q,d)$:

$$Score(q,d) = \sum_{t \in q \cap d} w_{t,d}$$

Walid Magdy, TTDS 2025/2026

THE UNIVERSITY of EDINBURGH

13

Document/Term vectors with tfidf

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	5.25	3.18	0	0	0	0.35
Brutus	1.21	6.1	0	1	0	0
Caesar	8.59	2.54	0	1.51	0.25	0
Calpurnia	0	1.54	0	0	0	0
Cleopatra	2.85	0	0	0	0	0
mercy	1.51	0	1.9	0.12	5.25	0.88
worser	1.37	0	0.11	4.15	0.25	1.95

→ Vector Space Model

Walid Magdy, TTDS 2025/2026

THE UNIVERSITY of EDINBURGH

14

7

原 PPT 第 7 页

人话直觉

好关键词要同时满足两件事：在这篇文档里声音大，在别的文档里不吵。

精确定义 / 机制：TF-IDF term weight 同时奖励 document 内高频和 collection 中稀有。它是经典强 baseline，简单、可解释、易实现。query-document score 可对共有 query terms 的权重求和，也可作为 VSM 中的向量权重。

$$w_{t,d} = (1 + \log_{10}tf(t,d)) \log_{10} \frac{N}{df(t)}$$


考试问法：若考 TF-IDF，必须定义 tf、df/idf，再解释“文档内多、集合中少”这一直觉。

7. Vector Space Model: 把文本放进几何空间

10/7/25

Vector Space Model

- Documents and Queries are presented as vectors
- Match (Q,D) = Distance between vectors
- Example: Q= Gossip Jealous
- Euclidean Distance?
Distance between the endpoints of the two vectors
- Large for vectors of diff. lengths
- Take a document d and append it to itself. Call this document d'.
 - "Semantically" d and d' have the same content
 - Euclidean distance can be quite large




Wolaid Maqdy, TTDS 2025/2026

THE UNIVERSITY of EDINBURGH

15

Angle Instead of Distance

- The angle between the two documents is 0, corresponding to maximal similarity.
- Key idea: Rank documents according to angle with query.
 - Rank documents in increasing order of the angle with query
 - Rank documents in decreasing order of cosine (query, document)
- Cosine of angle = projection of one vector on the other



Wolaid Maqdy, TTDS 2025/2026

THE UNIVERSITY of EDINBURGH

16

8

原 PPT 第 8 页

人话直觉

query 和 document 都变成箭头；相关性不是看箭头绝对多长，而是看它们指向是不是相近。

精确定义 / 机制：VSM 将 documents 和 queries 表示为同一 term space 中的 weighted vectors。每个维度是 term，值可由 TF-IDF 等 weighting 给出。匹配时比较 query vector 与 document vector 的相似度。

考试问法：答 VSM 时要说 shared term space、weighted vectors、similarity/ranking。

8. Length normalization: 别奖励“啰嗦”

10/7/25

Length Normalization

- A vector can be normalized by dividing each of its components by its length – for this we use the L_2 norm:

$$\|\vec{x}\|_2 = \sqrt{\sum_i x_i^2}$$

- Dividing a vector by its L_2 norm makes it a unit (length) vector (on surface of unit hypersphere)
- Effect on the two documents d and d' (d appended to itself) from earlier slide: they have identical vectors after length-normalization.
 - Long and short documents now have comparable weights

Walid Magdy, TTDS 2025/2026

THE UNIVERSITY OF EDINBURGH

17

Example

- $D1 = \begin{bmatrix} 1 \\ 3 \\ 2 \end{bmatrix} \rightarrow \|\vec{D1}\|_2 = \sqrt{1 + 9 + 4} = 3.74$
- $D1_{normalized} = \begin{bmatrix} 0.267 \\ 0.802 \\ 0.535 \end{bmatrix}$
- $D2 = \begin{bmatrix} 3 \\ 9 \\ 6 \end{bmatrix} \rightarrow \|\vec{D2}\|_2 = \sqrt{9 + 81 + 36} = 11.25$
- $D2_{normalized} = \begin{bmatrix} 0.267 \\ 0.802 \\ 0.535 \end{bmatrix}$

Walid Magdy, TTDS 2025/2026

THE UNIVERSITY OF EDINBURGH

18

9

原 PPT 第 9 页

人话直觉

一篇文档复制粘贴一遍，语义没变，但长度翻倍；如果直接按点积或欧氏距离，就会被长度骗。

精确定义 / 机制：长文档包含更多 terms，容易得到更高 raw score。L2 normalization 将向量除以其长度，使比较更关注方向而非规模。复制文档后向量长度变大，但 normalized direction 基本不变。

$$\|\vec{x}\|_2 = \sqrt{\sum_i x_i^2}, \quad \vec{x}' = \frac{\vec{x}}{\|\vec{x}\|_2}$$


考试问法：若问为什么不用 Euclidean distance，答 document length 会扭曲相似度；cosine/normalization 更稳。

9. Cosine similarity: 看夹角，不看体重

10/7/25

Cosine “Similarity” (Query, Document)

- $\vec{q}_i$ is the tf-idf weight of term i in the query
- $\vec{d}_i$ is the tf-idf weight of term i in the document
- For normalized vectors:
 $\cos(\vec{q}, \vec{d}) = \vec{q} \cdot \vec{d} = \sum_{i=1}^{|V|} q_i d_i$
- For non-normalized vectors:
$$\cos(\vec{q}, \vec{d}) = \frac{\vec{q} \cdot \vec{d}}{\|\vec{q}\| \|\vec{d}\|} = \frac{\vec{q}}{\|\vec{q}\|} \cdot \frac{\vec{d}}{\|\vec{d}\|} = \frac{\sum_{i=1}^{|V|} q_i d_i}{\sqrt{\sum_{i=1}^{|V|} q_i^2} \sqrt{\sum_{i=1}^{|V|} d_i^2}}$$



Walid Magdy, TTDS 2025/2026

THE UNIVERSITY OF EDINBURGH

19

Algorithm

```
COSINESCORE( $q$ )
1  float Scores[ $N$ ] = 0
2  float Length[ $N$ ]
3  for each query term  $t$ 
4  do calculate  $w_{t,q}$  and fetch postings list for  $t$ 
5    for each pair( $d, tf_{t,d}$ ) in postings list
6    do Scores[ $d$ ] +=  $w_{t,d} \times w_{t,q}$ 
7  Read the array Length
8  for each  $d$ 
9  do Scores[ $d$ ] = Scores[ $d$ ] / Length[ $d$ ]
10 return Top  $K$  components of Scores[]
```

Walid Magdy, TTDS 2025/2026

THE UNIVERSITY OF EDINBURGH

20

10

原 PPT 第 10 页

人话直觉

两个箭头方向越一致，说明 query 想要的词权重模式和 document 越接近。

精确定义 / 机制：cosine similarity 是 query vector 与 document vector 的点积除以各自长度。若 vectors 已归一化，cosine 就等于 dot product。它是 VSM 中最常见的 ranking score。

$$\cos(\vec{q}, \vec{d}) = \frac{\vec{q} \cdot \vec{d}}{\|\vec{q}\| \|\vec{d}\|}$$

考试问法：公式题要写 numerator 是 dot product，denominator 是两个 L2 norms；归一化后直接点积。

10. SMART notation: 给 weighting scheme 起缩写名

10/7/25

TFIDF Variants

Term frequency	Document frequency	Normalization
n (natural) $tf_{t,d}$	n (no) 1	n (none) 1
l (logarithm) $1 + \log(tf_{t,d})$	t (idf) $\log \frac{N}{df_t}$	c (cosine) $\frac{1}{\sqrt{w_q^2 + w_d^2}}$
a (augmented) $0.5 + \frac{0.5 \cdot tf_{t,d}}{\max(tf_{t,d})}$	p (prob idf) $\max(0, \log \frac{N - df_t + 1}{df_t})$	u (pivoted unique) $1/u$
b (boolean) $\begin{cases} 1 & \text{if } tf_{t,d} > 0 \\ 0 & \text{otherwise} \end{cases}$		b (byte size) $1/CharLength^\alpha$, $\alpha < 1$
L (log ave) $\frac{1 + \log(tf_{t,d})}{1 + \log(\max_c(tf_{c,d}))}$		

- Many search engines allow for different weightings for queries vs. documents
- **SMART** Notation: use notation *ddd.qqq*, using the acronyms from the table
- A very standard weighting scheme is: *Inc.Itc*

Walid Magdy, TTDS 2025/2026 THE UNIVERSITY OF EDINBURGH

21

For Lab and CW

Term frequency	Document frequency	Normalization
n (natural) $tf_{t,d}$	n (no) 1	n (none) 1
l (logarithm) $1 + \log(tf_{t,d})$	t (idf) $\log \frac{N}{df_t}$	c (cosine) $\frac{1}{\sqrt{w_q^2 + w_d^2}}$
a (augmented) $0.5 + \frac{0.5 \cdot tf_{t,d}}{\max(tf_{t,d})}$	p (prob idf) $\max(0, \log \frac{N - df_t + 1}{df_t})$	u (pivoted unique) $1/u$
b (boolean) $\begin{cases} 1 & \text{if } tf_{t,d} > 0 \\ 0 & \text{otherwise} \end{cases}$		b (byte size) $1/CharLength^\alpha$, $\alpha < 1$
L (log ave) $\frac{1 + \log(tf_{t,d})}{1 + \log(\max_c(tf_{c,d}))}$		

“OR” operator, then:

$$Score(q, d) = \sum_{t \in q \cap d} (1 + \log_{10} tf(t, d)) \times \log_{10} \left(\frac{N}{df(t)} \right)$$

Walid Magdy, TTDS 2025/2026 THE UNIVERSITY OF EDINBURGH

22

11

原 PPT 第 11 页

人话直觉

SMART 像配方标签：document 怎么处理、query 怎么处理，用六个字母浓缩写清楚。

精确定义 / 机制：SMART notation 用 *ddd.qqq* 描述 document 和 query 的 term weighting scheme，包括 tf scaling、idf usage、normalization 等。例如 *Inc.Itc* 表示 documents 和 queries 使用不同组合。它帮助系统化比较 weighting variants。

考试问法：若问 SMART，不必背所有组合细节，但要知道它描述 query/document 的 weighting scheme。

考试速记版

- Ranked retrieval 给 document 打 score 并返回 top-k；Boolean 只有 match/no match。
- Jaccard 是 |交集|/|并集|，缺点是不看 tf、idf 和长度。
- TF 表示词在文档中出现次数；log TF 体现边际收益递减。

- IDF 表示词的集合稀有性：df 越大，idf 越小。
- TF-IDF = 文档内高频 × 集合中稀有，是强 baseline。
- VSM 把 query/document 表示为 term vectors，用 cosine similarity 排序。
- Length normalization 防止长文档仅因更长而占便宜。

一段稳妥考场回答

Ranked retrieval 解决 Boolean retrieval 只有 match/no match、不能排序的问题。系统为每个 document 对 query 计算 score，并返回 top-k。最简单的 Jaccard 只看集合 overlap，但不考虑 term frequency、term rarity 和 length normalization。TF-IDF 用 tf 衡量 term 在文档中的重要性，用 idf 衡量它在 collection 中的区分能力，常写为 $(1+\log tf)\log(N/df)$ 。在 Vector Space Model 中，query 和 document 都表示为 TF-IDF weighted vectors，再用 cosine similarity 比较方向相似度；L2 normalization 可减少文档长度差异造成的偏置。